

LoRA-Diffusion: Parameter-Efficient Fine-Tuning via Low-Rank Trajectory Decomposition

Iman Khazrak

Department of Computer Science
Bowling Green State University
ikhazra@bgsu.edu

Robert Green

Department of Computer Science
Bowling Green State University
greenr@bgsu.edu

February 8, 2026

Abstract

Parameter-efficient fine-tuning methods such as LoRA have transformed the adaptation of large autoregressive language models, enabling task-specific customization with fewer than 1% trainable parameters. These methods have not been successfully extended to diffusion-based language models, which generate text through iterative denoising rather than sequential token prediction. We propose LoRA-Diffusion, a parameter-efficient fine-tuning approach that applies low-rank decomposition to the denoising trajectory instead of model weights. Unlike weight-based LoRA, which modifies individual transformation matrices, our method learns low-rank perturbations to the entire diffusion path from noise to output. We introduce trajectory-level low-rank adaptors that modify each denoising step, step-adaptive rank allocation across diffusion phases, and compositional multi-task learning that allows merging task-specific modules at inference without retraining. On SST-2, QNLI, and MRPC (5 seeds) we report token-level denoising validation accuracy; LoRA-Diffusion reaches the highest mean on SST-2 (88.01%) and strong performance on QNLI (99.39%) and MRPC (97.56%). Joint multi-task training (same 5 methods, 5 seeds) shows LoRA-Diffusion best ($96.88\% \pm 0.44\%$ token-level). LoRA-Diffusion achieves the highest SST-2 validation accuracy while training 28.7% of parameters (instruction encoder 27.5% + trajectory adapters 1.2%)¹. The approach is competitive with adapter layers and baselines, and reduces storage per task (151 MB vs. 525 MB for full fine-tuning) while exhibiting minimal catastrophic forgetting. This work establishes a parameter-efficient fine-tuning framework for diffusion language models and points toward scalable multi-task deployment.

1 Introduction

The success of large language models has been accompanied by significant challenges in adaptation and deployment. Full fine-tuning of billion-parameter models is computationally costly, requiring substantial GPU memory and training time Brown et al. [2020]. Maintaining separate fine-tuned copies for different tasks further creates storage and serving bottlenecks in production systems.

Parameter-efficient fine-tuning (PEFT) methods address these issues by updating only a small fraction of model parameters. Among them, Low-Rank Adaptation (LoRA) has proven especially effective, achieving near-full fine-tuning performance on autoregressive models while training fewer than 1% of parameters Hu et al. [2021]. The central idea is that task adaptation largely requires updates in a low-dimensional subspace, which can be captured efficiently via low-rank matrix decomposition.

¹Parameter accounting: total trainable 39.6M = instruction encoder 37.8M + trajectory adapters 1.7M; percentages relative to base model 137.7M.

Recent work has shown that discrete diffusion models can match or exceed autoregressive models in text generation quality Lou et al. [2023], Sahoo et al. [2024]. Diffusion models offer bidirectional context, parallel generation, controllable generation, and diverse sampling. Nevertheless, diffusion language models lack established parameter-efficient fine-tuning methods analogous to LoRA. Existing approaches either apply standard LoRA to diffusion weights (treating the model as a standard transformer), perform full fine-tuning, or use adapter layers or prefix tuning, which introduce sequential bottlenecks. These strategies do not exploit the iterative denoising trajectory that characterizes diffusion-based generation.

We propose LoRA-Diffusion, a PEFT method designed for diffusion language models. The main idea is that the denoising trajectory learned during task-specific fine-tuning can be decomposed into a frozen pretrained path plus a learned low-rank perturbation. Formally, we write

$$\mathbf{x}_t^{\text{fine-tuned}} = \mathbf{x}_t^{\text{pretrained}} + \Delta \mathbf{x}_t^{\text{low-rank}}, \quad (1)$$

where $\Delta \mathbf{x}_t^{\text{low-rank}}$ is produced by lightweight low-rank adaptors conditioned on the task instruction. The perturbation is applied in *hidden representation space*: $\mathbf{h}'_t = \mathbf{h}_t + \delta_t$, then logits $\mathbf{l}_t = \text{OutputHead}(\mathbf{h}'_t)$ and \mathbf{x}_{t-1} from the output head (Section 3, Eq. 2 and Eq. 3). Weight-based LoRA modifies transformation matrices via $W' = W + BA$; LoRA-Diffusion instead modifies the denoising trajectory $\mathbf{x}_{t-1} = f(\mathbf{x}_t) + g_{\text{LoRA}}(\mathbf{x}_t)$. Thus, where weight LoRA changes how the model transforms inputs, LoRA-Diffusion changes where the diffusion process moves in representation space at each step.

We make the following contributions. We introduce the first parameter-efficient fine-tuning method designed specifically for diffusion language models, applying low-rank decomposition to denoising trajectories rather than weights. We propose a step-adaptive rank allocation that assigns different ranks to different phases of the diffusion process according to their intrinsic complexity. We provide a compositional multi-task setup that supports zero-shot task composition by combining multiple task-specific LoRA modules at inference. We present an empirical evaluation on SST-2, QNLI, and MRPC (single-task and joint multi-task) with a BERT-based diffusion model (137.7M parameters), 5 seeds (42–46), comparing LoRA-Diffusion to full fine-tuning and several PEFT baselines (weight LoRA, adapters, BitFit), with token-level denoising accuracy, efficiency metrics (trainable parameters, storage, training time, inference latency), and ablations for rank and orthogonality regularization. We give an information-theoretic motivation for trajectory-level low-rank structure and clarify positioning versus adapter layers and timestep-aware weight LoRA (T-LoRA, FouRA). We release an open-source implementation to support reproducibility and extension.

The rest of the paper is organized as follows. Section 2 reviews related work on diffusion models for language, parameter-efficient fine-tuning, and multi-task learning. Section 3 presents our methodology, including preliminaries, trajectory-level low-rank adaptation, the training objective, multi-task composition, and implementation details. Section 4 describes the experimental setup and results on SST-2, QNLI, and MRPC, including single-task and multi-task GLUE results, main results, efficiency analysis, catastrophic forgetting, ablations, and comparison with weight-based LoRA. Section 5 summarizes our contributions, discusses limitations and future work, and closes with broader impact and reproducibility notes.

2 Related Work

2.1 Diffusion Models for Language

Austin et al. [2021] introduced discrete diffusion for categorical data, with uniform and absorbing-state transition mechanisms. Hoogeboom et al. [2021] proposed argmax flows for multinomial diffusion. More recently, Lou et al. [2023] presented SEDD, which achieves competitive generation quality with autoregressive models; Sahoo et al. [2024] simplified the setup with masked

diffusion; and Li et al. [2022] explored controlled generation with Diffusion-LM. All of this work focuses on pretraining or basic fine-tuning. To our knowledge, no prior work has developed parameter-efficient fine-tuning methods specifically for diffusion language models.

2.2 Parameter-Efficient Fine-Tuning

Hu et al. [2021] introduced LoRA for low-rank adaptation of autoregressive models. Dettmers et al. [2023] combined LoRA with quantization (QLoRA), and Zhang et al. [2023] proposed AdaLoRA to adapt ranks dynamically. Other PEFT methods include prefix tuning Li et al. [2021], prompt tuning Lester et al. [2021], adapter layers Houlsby et al. [2019], and BitFit Zaken et al. [2021], which trains only bias terms. These methods target autoregressive architectures. Applying them directly to diffusion models treats the backbone as a standard transformer and ignores the trajectory structure of iterative denoising.

Recent work has explored timestep-aware and rank-adaptive PEFT for diffusion models, primarily in the image domain. T-LoRA [2024] (T-LoRA) applies timestep-dependent rank masking and orthogonalization to maintain effective rank across diffusion steps. FouRA [2024] (FouRA) introduces frequency-domain LoRA with adaptive rank gating across timesteps. TALoRA [2024] (TALoRA) and MSFP [2024] (MSFP) propose timestep-adaptive low-rank factorization with hub-based sharing. SeLoRA [2024] (SeLoRA) and GeLoRA [2024] (GeLoRA) provide principled rank allocation based on Fisher information and intrinsic dimension. EST-LoRA [2024] (EST-LoRA) studies training-free adapter fusion via routing at inference. TC-LoRA [2024] (TC-LoRA) conditions low-rank weight updates on timestep and condition via a hypernetwork, modulating weight functions per timestep/condition. EfficientDM [2023] (EfficientDM) and Glance [2024] (Glance) demonstrate practical PEFT/acceleration strategies with step/phase specializations. Delta Sampling [2024] (Delta Sampling) operates at inference by reusing deltas in prediction space. These methods operate in weight or frequency space and allocate capacity across timesteps, but do not explicitly model trajectory-level perturbations. LoRA-Diffusion differs by operating directly in representation/trajecotry space, where low-rank structure emerges naturally from the iterative denoising process, and by using a phase-shared design that keeps parameter counts independent of the number of diffusion steps. Unlike TC-LoRA which modulates weights, LoRA-Diffusion modulates trajectory corrections, offering different representational advantages and computational costs.

Comparison with timestep-aware and diffusion PEFT. Table 1 summarizes how LoRA-Diffusion relates to prior PEFT methods. Key trade-offs: (1) **Compute:** LoRA-Diffusion adds a lightweight g_ϕ per diffusion step, so inference cost is higher than weight LoRA unless g_ϕ is very small; (2) **Compositionality:** trajectory superposition (router-weighted sum of task adapters) vs. weight-space task arithmetic; (3) **Trainability:** trajectory-only adapters are 1.2% of base; with instruction encoder, total trainable is 28.7%.

2.3 Multi-Task Learning and Low-Rank Theory

Ilharco et al. [2022] showed that task vectors can be combined via task arithmetic. Wang et al. [2020] used orthogonal subspace projection to reduce interference. Routing-based mixture-of-experts approaches Fedus et al. [2022] select experts per input. Aghajanyan et al. [2020] demonstrated that task adaptation has low intrinsic dimensionality; Li et al. [2018] measured intrinsic dimensionality empirically. Tishby and Zaslavsky [2015] provided an information-theoretic perspective via the information bottleneck. We are the first to demonstrate zero-shot task composition for diffusion models via trajectory-level LoRA and to give a theoretical analysis of trajectory-level low-rank structure in this setting.

Table 1: Comparison with diffusion and timestep-aware PEFT.

Method	What is adapted	Timestep-aware	Composition	Domain
Full Fine-Tuning	Weights	—	—	Text/Image
BitFit	Biases	No	—	Text
Prefix Tuning	Prompts	No	Limited	Text
Adapters	Activations (layer)	No	Task arithmetic	Text
LORA (weight)	Weights W	No	Task arithmetic	Text/Image
T-LoRA	Weights (rank mask)	Yes	—	Image
FouRA	Weights (frequency)	Yes	—	Image
SeLoRA / GeLoRA	Weights (rank alloc.)	Yes	—	Image
EST-LoRA	Weights (routing)	Yes	Routing	Image
Delta Sampling	Predictions (inference)	Yes	—	Text
LoRA-Diffusion (Ours)	Trajectory \mathbf{h}_t	Yes	Router / superposition	Text

3 Methodology

3.1 Preliminaries

A discrete diffusion model for language defines a forward Markov process that gradually corrupts clean text $\mathbf{x}_0 = (x_0^1, \dots, x_0^n)$, $x_0^i \in \mathcal{V}$, over timesteps $t \in [1, T]$. Common transitions include the uniform and absorbing-state (masking) schemes of Austin et al. [2021]. The model learns to reverse the process by predicting \mathbf{x}_0 from \mathbf{x}_t and t , and is trained with a simplified objective $\mathcal{L}_{\text{simple}} = \mathbb{E}_{\mathbf{x}_0, t, \mathbf{x}_t}[-\log p_\theta(\mathbf{x}_0 \mid \mathbf{x}_t, t)]$. For conditional generation, conditioning c (e.g. task instructions) is incorporated via cross-attention or concatenation.

LoRA Hu et al. [2021] adapts pretrained weights W_0 via $W = W_0 + BA$, with $B \in \mathbb{R}^{d \times r}$, $A \in \mathbb{R}^{r \times d}$, $r \ll d$, and only B and A trained. Its success is tied to the low intrinsic dimensionality of task adaptation Aghajanyan et al. [2020]. Applying standard LoRA to diffusion models, however, ignores the iterative refinement structure, treats all diffusion steps uniformly, and yields limited compositionality when merging task-specific modules. We therefore move from weight-level to trajectory-level adaptation.

3.2 Representation Space and Trajectory Perturbations

In discrete diffusion language models, \mathbf{x}_t represents discrete token IDs from the vocabulary \mathcal{V} . The model operates on hidden representations $\mathbf{h}_t = \text{Transformer}(\mathbf{x}_t, t)$ obtained by passing token embeddings through the transformer backbone with time embeddings. The output head then computes logits $\mathbf{l}_t = \text{OutputHead}(\mathbf{h}_t)$ to predict the next token distribution.

Trajectory perturbations are applied in the hidden representation space, not directly to tokens or logits. The data flow is: **tokens** \mathbf{x}_t (discrete IDs) \rightarrow **embeddings** \rightarrow **hidden states** $\mathbf{h}_t = \text{Transformer}(\mathbf{x}_t, t) \rightarrow$ **perturbation** $\mathbf{h}'_t = \mathbf{h}_t + \delta_t \rightarrow$ **logits** $\mathbf{l}_t = \text{OutputHead}(\mathbf{h}'_t)$. Specifically:

$$\mathbf{h}'_t = \mathbf{h}_t + \delta_t, \quad (2)$$

where δ_t is the learned low-rank perturbation, and then $\mathbf{l}_t = \text{OutputHead}(\mathbf{h}'_t)$. This preserves the probabilistic structure because: (1) the output head remains deterministic, (2) perturbations are learned to maintain valid conditional distributions $p(\mathbf{x}_0 \mid \mathbf{x}_t, t, c)$, and (3) the training objective ensures the perturbed trajectory produces valid reverse diffusion transitions. The output head is deterministic, so $p(\mathbf{x}_0 \mid \mathbf{x}_t, t, c)$ stays well-defined; the denoising loss trains δ_t to yield valid reverse transitions.

Algorithm. At each denoising step t , the pipeline is: (1) Compute $\mathbf{h}_t = \text{Transformer}(\mathbf{x}_t, t)$; (2) Compute $\delta_t = \sum_i \sigma(t) \cdot g_{\phi_i}(\mathbf{h}_t, t, c)$; (3) Set $\mathbf{h}'_t = \mathbf{h}_t + \delta_t$; (4) Compute $\mathbf{l}_t = \text{OutputHead}(\mathbf{h}'_t)$;

(5) Obtain \mathbf{x}_{t-1} from \mathbf{l}_t (e.g., sampling or argmax). Thus perturbations modify *hidden states* \mathbf{h}_t , not token IDs or logits directly. Implementation: training in `src/training/losses.py` (LoRA perturbation applied to hidden states, then output head); inference in `src/models/base_diffusion.py` (`sample_with_composition`); LoRA modules in `src/models/lora_modules.py` (`TrajectoryLoRAAdapter`, `LoRADiffusionModule`).

3.3 Trajectory-Level Low-Rank Adaptation

At each denoising step t , the model computes $\mathbf{x}_{t-1} = f_\theta(\mathbf{x}_t, t, c)$ via the process: tokens $\mathbf{x}_t \rightarrow$ hidden states $\mathbf{h}_t \rightarrow$ (optionally perturbed) $\mathbf{h}'_t \rightarrow$ logits $\mathbf{l}_t \rightarrow$ predicted tokens \mathbf{x}_{t-1} . After task-specific fine-tuning, the denoising function changes from f_θ to $f_{\theta'}$. We hypothesize that the difference $\Delta f = f_{\theta'} - f_\theta$ can be well approximated by a low-rank function in representation space, i.e. that the trajectory perturbation δ_t lies in a low-dimensional subspace of \mathbb{R}^d where d is the hidden dimension.

We decompose the fine-tuned trajectory as

$$\mathbf{x}_{t-1}^{\text{fine-tuned}} = \underbrace{f_{\theta_0}(\mathbf{x}_t, t, c)}_{\text{frozen pretrained}} + \underbrace{\sum_{i=1}^k \sigma(t) \cdot g_{\phi_i}(\mathbf{x}_t, t, c)}_{\text{learnable low-rank perturbation}}, \quad (3)$$

where f_{θ_0} is the frozen pretrained denoising function, g_{ϕ_i} is the i -th low-rank perturbation module, $\sigma(t)$ is a step-adaptive scaling function, and k is the number of LoRA modules per step (typically 1–4).

Each module g_{ϕ_i} is implemented as $g_{\phi_i}(\mathbf{x}_t, t, c) = A_i(c) \cdot \text{ReLU}(B_i(\mathbf{x}_t, t))$, with $B_i : \mathbb{R}^d \rightarrow \mathbb{R}^r$ (down-projection) and $A_i : \mathbb{R}^r \rightarrow \mathbb{R}^d$ (up-projection), $r \ll d$. The down-projection is $B_i(\mathbf{x}_t, t) = W_B^{(i)}[\mathbf{x}_t; \text{Emb}(t)]$, while the up-projection is implemented via FiLM-style conditioning: a base matrix plus instruction-dependent scale and shift. Concretely, we realize $A_i(c)$ as

$$A_i(c)v = W_A^{(i)}(\gamma_i(c) \odot v) + \beta_i(c), \quad (4)$$

where $\gamma_i(c)$ and $\beta_i(c)$ are computed by a lightweight instruction encoder and \odot denotes elementwise multiplication. The **nominal rank** r refers to the bottleneck dimension of $B_i(\cdot, t)$; $A_i(c)$ is a conditional up-projection. FiLM applies elementwise scale and shift to the bottleneck vector $v \in \mathbb{R}^r$; the output remains in the column space of $W_A^{(i)}$ (plus a fixed shift per c), so for each fixed c the map $v \mapsto A_i(c)B_i(\mathbf{h}_t, t)$ has range in an at-most- r -dimensional affine subspace and **effective rank at most** r . We measure effective rank empirically (Section 4); the script `scripts/analyze_effective_rank.py` computes singular value spectra and effective rank per phase. The nuclear norm $\mathcal{R}_{\text{rank}}$ is applied to the base matrices $W_A^{(i)}, W_B^{(i)}$, encouraging low-rank structure in the unconstrained components.

Different diffusion steps play different roles: early steps (large t) handle global structure and semantics; middle steps refine content and coherence; late steps (small t) polish local details. We partition timesteps into three phases: **Early** ($t > 2T/3$), **Mid** ($T/3 < t \leq 2T/3$), and **Late** ($t \leq T/3$). For $T = 100$, this corresponds to early: $t \in [67, 100]$, mid: $t \in [34, 66]$, and late: $t \in [0, 33]$. We use step-adaptive scaling $\sigma(t)$ with $\sigma_{\text{early}} = 1.0$, $\sigma_{\text{mid}} = 0.5$, and $\sigma_{\text{late}} = 0.25$. We also allocate rank $r(t)$ adaptively: $r_{\text{early}} = 64$, $r_{\text{mid}} = 32$, and $r_{\text{late}} = 8$. Early steps explore a high-dimensional space of global structures and thus use higher rank; late steps refine within a local neighborhood and use lower rank. In the reference implementation we instantiate three banks of adapters corresponding to early/mid/late phases and reuse them across all timesteps within a phase, so the trainable parameter count is independent of T and step-awareness is expressed through the phase-dependent scaling $\sigma(t)$ rather than separate parameters for every timestep.

3.4 Training Objective

The training objective is

$$\mathcal{L} = \mathcal{L}_{\text{denoise}} + \lambda_{\text{rank}} \mathcal{R}_{\text{rank}} + \lambda_{\text{orth}} \mathcal{R}_{\text{orth}}, \quad (5)$$

with $\mathcal{L}_{\text{denoise}} = \mathbb{E}_{\mathbf{x}_0, c, t, \mathbf{x}_t}[-\log p_\theta(\mathbf{x}_0 \mid \mathbf{x}_t, t, c)]$ and

$$\mathcal{R}_{\text{rank}} = \sum_{i=1}^k \|W_A^{(i)}\|_* + \|W_B^{(i)}\|_*, \quad (6)$$

$$\mathcal{R}_{\text{orth}} = \sum_{i \neq j} \|W_A^{(i)T} W_A^{(j)}\|_F^2. \quad (7)$$

The nuclear norm encourages low-rank structure; the orthogonality term encourages complementary learned directions. We use $\lambda_{\text{rank}} = 0.01$, $\lambda_{\text{orth}} = 0.001$, learning rate 1×10^{-4} for LoRA parameters only, and keep the base model frozen. Regularization ablation is reported in the supplement.

3.5 Multi-Task Composition

For each task j , we train a separate set of LoRA modules $\{\phi_i^{(j)}\}$. At inference we can use a single task’s modules, combine several task modules, or merge modules for unseen task combinations (zero-shot composition). Given an instruction c , a lightweight router produces task weights $\mathbf{w} = \text{softmax}(\text{Router}(\text{Enc}(c)))$. The composed update is

$$\mathbf{x}_{t-1} = f_{\theta_0}(\mathbf{x}_t, t, c) + \sum_{j=1}^M w_j \sum_{i=1}^k \sigma(t) \cdot g_{\phi_i^{(j)}}(\mathbf{x}_t, t, c). \quad (8)$$

The router is a 2-layer MLP with 512 hidden units and $\sim 1\text{M}$ parameters, trained jointly with the LoRA modules via multi-task learning.

3.6 Inference Procedure

Algorithm 1 summarizes inference. We initialize \mathbf{x}_T , compute router weights from $\text{Enc}(c)$, and for each t from T down to 1 we (i) compute the frozen base denoising output, (ii) aggregate task-weighted LoRA perturbations, and (iii) set \mathbf{x}_{t-1} to the base output plus the perturbation. We return \mathbf{x}_0 .

Algorithm 1 LoRA-Diffusion Inference

```
1: Input: Instruction  $c$ , diffusion steps  $T$ , LoRA modules  $\{\phi_i^{(j)}\}_{j=1}^M$ 
2: Initialize:  $\mathbf{x}_T \sim \text{Uniform}(\mathcal{V})$  or  $\mathcal{N}(0, I)$  (depending on forward process)
3:  $\mathbf{w} \leftarrow \text{Router}(\text{Enc}(c))$ 
4:  $t \leftarrow T$ 
5: while  $t \geq 1$  do
6:    $\mathbf{x}_t^{\text{base}} \leftarrow f_{\theta_0}(\mathbf{x}_t, t, c)$ 
7:    $\boldsymbol{\delta} \leftarrow \mathbf{0}$ 
8:    $j \leftarrow 1$ 
9:   while  $j \leq M$  do
10:     $i \leftarrow 1$ 
11:    while  $i \leq k$  do
12:       $\boldsymbol{\delta} \leftarrow \boldsymbol{\delta} + w_j \cdot \sigma(t) \cdot g_{\phi_i^{(j)}}(\mathbf{x}_t, t, c)$ 
13:       $i \leftarrow i + 1$ 
14:    end while
15:     $j \leftarrow j + 1$ 
16:  end while
17:   $\mathbf{x}_{t-1} \leftarrow \mathbf{x}_t^{\text{base}} + \boldsymbol{\delta}$ 
18:   $t \leftarrow t - 1$ 
19: end while
20: Return  $\mathbf{x}_0$ 
```

3.7 Implementation Details

We use SEDD Lou et al. [2023] as the base diffusion model. Table 2 gives model configurations. Table 3 lists LoRA-Diffusion hyperparameters. For our BERT setup ($d = 768$, $T = 100$, $k = 2$), the total trainable parameters are 39.6M (28.7% of base model 137.7M), including the instruction encoder (37.8M, 27.5%) and trajectory adapters (1.7M, 1.2%). The phase-shared design keeps parameter counts independent of T . Table 6 gives a single, consistent accounting for all methods.

Table 2: Base model configurations (illustrative; our experiments use a BERT-based model with 137.7M trainable parameters, corresponding roughly to the Small configuration).

Configuration	Small	Medium	Large
Parameters	350M	1.3B	7B
Layers	12	24	32
Hidden dimension	1024	2048	4096
Attention heads	16	32	32
FFN dimension	4096	8192	16384
Vocabulary size	50k	50k	50k
Max sequence length	512	1024	2048
Diffusion steps T	100	100	100

3.8 Theoretical Justification

Under the information bottleneck principle Tishby et al. [2000], task adaptation learns a compressed representation $\mathbf{z}_{\text{task}} \in \mathbb{R}^r$. If the trajectory perturbation $\Delta \mathbf{x}_t$ lies approximately in an r -dimensional subspace, it can be written as $\Delta \mathbf{x}_t = A \mathbf{z}_{\text{task}} + \epsilon$ with small ϵ , which matches the low-rank structure used by LoRA-Diffusion. We define the effective rank of trajectory perturbations via the entropy of normalized singular values; empirically, $r_{\text{eff}} \ll d$ across steps, and early

Table 3: LoRA-Diffusion hyperparameters.

Hyperparameter	Value
Rank (early, $t > 2T/3$)	64
Rank (middle, $T/3 < t \leq 2T/3$)	32
Rank (late, $t \leq T/3$)	8
Number of LoRA modules k	2
Scaling $\sigma_{\text{high}}, \sigma_{\text{mid}}, \sigma_{\text{low}}$	1.0, 0.5, 0.25
$\lambda_{\text{rank}}, \lambda_{\text{orth}}$	0.01, 0.001
Learning rate	1×10^{-4}
Batch size	64 (with gradient accumulation)
Training steps	10k–20k (task-dependent)

steps exhibit higher effective rank than late steps, consistent with our step-adaptive allocation. A script `analyze_effective_rank.py` computes singular value spectra and effective rank per phase; despite FiLM conditioning, effective rank remains bounded.

Table 4 compares PEFT methods. Table 5 contrasts weight LoRA with LoRA-Diffusion. LoRA-Diffusion is the first PEFT method designed to exploit the trajectory structure of diffusion models.

Table 4: Comparison of parameter-efficient fine-tuning methods.

Method	Key idea	Trainable %	Compatible with diffusion?
Full Fine-Tuning	Update all parameters	100%	Yes
BitFit	Train only bias terms	0.1%	Partially
Prefix Tuning	Prepend learnable prompts	0.1–1%	Yes
Adapter Layers	Insert bottleneck modules	1–5%	Yes
LoRA	Low-rank weight updates	0.1–1%	Naive application
LoRA-Diffusion (Ours)	Low-rank trajectory updates	28.7% (1.2% adapters only)	Designed for

Table 5: Conceptual comparison: Weight LoRA vs. LoRA-Diffusion.

Aspect	Weight LoRA	LoRA-Diffusion
What is decomposed?	Matrices W	Trajectories $\mathbf{x}_t \rightarrow \mathbf{x}_{t-1}$
Where is low-rank applied?	Parameter space	Representation space
Frozen component	W_0	f_{θ_0}
Learned component	$\Delta W = BA$	$\Delta \mathbf{x}_t = g_{\phi}(\mathbf{x}_t)$
Compositionality	Limited (interference)	Natural (superposition)
Step-awareness	No	Yes (adaptive rank)

4 Experiments and Results

4.1 Experimental Setup

We evaluate on the SST-2 sentiment classification task with a base model architecture based on SEDD Lou et al. [2023]. The model uses a BERT-based transformer backbone with 137.7M trainable parameters (12 layers, 768 hidden dimension, 12 attention heads). **Full fine-tuning** updates all 137.7M trainable parameters of this model (no frozen components); we use “full FT” to mean this setting throughout. We report **validation accuracy using the same metric**

as training: token-level denoising accuracy (fraction of masked tokens predicted correctly on the validation set). We do not use generation or a separate classification head. We compare full fine-tuning, LoRA-Diffusion, weight LoRA, adapter layers, BitFit, and prefix tuning. Weight LoRA uses rank 64 on Q , K , V , O , and MLP layers; prefix tuning uses length 32; adapters use bottleneck dimension 256. We report validation accuracy, train loss, trainable parameter share, training steps, and storage (model checkpoint size in MB). Experiments use $4 \times$ NVIDIA A100 40GB GPUs, PyTorch 2.0, Hugging Face Transformers, AdamW with learning rate 1×10^{-4} and cosine decay, 500 warmup steps, effective batch size 64 with gradient accumulation, and FP16 mixed precision. We tune learning rate and regularization on the validation set.

Statistical rigor and reproducibility. For GLUE single- and multi-task experiments we use 5 random seeds (42–46). For each method-task combination, we report mean \pm standard deviation across seeds. We use paired t-tests to assess statistical significance between methods (Table 17). Timing and latency in the appendix use 10 seeds (42–51) from separate runs. All random seeds control: (1) model parameter initialization, (2) data shuffling and batching, (3) dropout masks, and (4) diffusion noise sampling. We set `torch.manual_seed`, `np.random.seed`, and `random.seed` for full reproducibility.

4.2 Statistical Analysis

We employ rigorous statistical methods to validate our findings:

Descriptive statistics. For each metric, we report mean (μ), standard deviation (σ), and 95% confidence interval (CI). The CI is computed as $\mu \pm t_{0.975,n-1} \cdot \text{SEM}$, where $\text{SEM} = \sigma/\sqrt{n}$ is the standard error of the mean; for GLUE we use $n = 5$ seeds.

Significance testing. We use paired t-tests to compare methods, treating each seed as a paired observation. For method A vs. method B , we test the null hypothesis $H_0 : \mu_A = \mu_B$ against the alternative $H_1 : \mu_A \neq \mu_B$. We report two-tailed p-values and apply Bonferroni correction when making multiple comparisons (e.g., comparing LoRA-Diffusion against 4 baselines: $\alpha_{\text{corrected}} = 0.05/4 = 0.0125$).

Effect size. We compute Cohen’s d to quantify practical significance: $d = (\mu_A - \mu_B)/\sigma_{\text{pooled}}$, where $\sigma_{\text{pooled}} = \sqrt{(\sigma_A^2 + \sigma_B^2)/2}$. We interpret $|d| < 0.2$ as negligible, $0.2 \leq |d| < 0.5$ as small, $0.5 \leq |d| < 0.8$ as medium, and $|d| \geq 0.8$ as large effect.

Robustness checks. We verify assumptions: (1) normality via Shapiro-Wilk test, (2) homogeneity of variance via Levene’s test. When assumptions are violated, we use non-parametric Wilcoxon signed-rank test as a robustness check.

SST-2 Formulation: For SST-2 sentiment classification, we format each example as an instruction-following task. The input sentence is embedded in an instruction template; the model is conditioned on this instruction. All reported validation and test accuracies are **token-level denoising accuracy**: the fraction of masked tokens predicted correctly on the validation/test split, using the same metric as training.

Parameter and storage accounting. We report a single, consistent accounting for all methods. Base model: 137.7M trainable parameters. LoRA-Diffusion total: 39.6M (28.7%), comprising instruction encoder 37.8M (27.5%) and trajectory adapters 1.7M (1.2%). All percentages are relative to the base model 137.7M. Table 6 summarizes trainable parameters and storage for each method.

Why 28.7% for LoRA-Diffusion vs. less for others? LoRA-Diffusion’s 28.7% is not a tuned budget but the outcome of its design: the method conditions on the task instruction via a shared *instruction encoder* (27.5%) that maps instructions to adapter weights, plus lightweight *trajectory adapters* (1.2%) that apply low-rank updates to the denoising path. The baselines (weight LoRA, adapters, BitFit, prefix tuning) are applied in their standard configurations from the literature (e.g., rank 64 for weight LoRA, bottleneck 256 for adapters), which yield lower trainable fractions. We compare each method in its natural or standard configuration rather

than at a fixed parameter budget; this reflects how each approach would typically be used. Despite using more parameters than other PEFT methods, LoRA-Diffusion remains much more parameter-efficient than full fine-tuning (28.7% vs. 100%) and achieves the highest validation accuracy in our experiments, so the comparison is meaningful both for efficiency (vs. full FT) and for accuracy (vs. all baselines).

Table 6: Parameter and storage accounting (base model 137.7M). All PEFT methods: trainable parameters only; full fine-tuning: full model.

Method	Trainable params	% of base	Storage (MB)
Full Fine-Tuning	137.7M	100.0%	525.2
LoRA-Diffusion (total)	39.6M	28.7%	150.9
Instruction encoder	37.8M	27.5%	—
Trajectory adapters only	1.7M	1.2%	—
Weight LoRA	9.7M	6.6%	36.9
Adapters	18.9M	12.1%	72.2
BitFit	156.5K	0.1%	0.6
Prefix Tuning	9.9M	7.2%	37.7

4.3 Main Results

Table 7 reports performance versus trainable parameters. We report **training accuracy** and **validation accuracy** using the **same metric**: token-level denoising accuracy (fraction of masked tokens predicted correctly on the training set and on the validation set, respectively). LoRA-Diffusion uses 28.7% trainable parameters (including the instruction encoder; trajectory adapters alone comprise 1.2%). Relative performance (Val acc. as % of full fine-tuning) is the primary comparison. Prefix tuning was not fully implemented in our diffusion setup.

Table 7: Performance on SST-2 (mean \pm std over 5 seeds, seeds 42–46, from `outputs/glue_single/`). Train acc. and Val acc. = token-level denoising (same metric). Relative = Val acc. as % of full fine-tuning.

Method	Trainable %	Train acc. (%)	Val acc. (%)	Relative
Full Fine-Tuning	100.0	85.52 ± 0.72	84.81 ± 0.38	100.0%
LoRA-Diffusion	28.7	87.99 ± 0.54	88.01 ± 0.27	103.8%
Weight LoRA	6.6	85.23 ± 0.72	85.23 ± 0.32	100.5%
Adapter Layers	12.1	85.20 ± 0.61	85.17 ± 0.07	100.4%
BitFit	0.1	85.05 ± 0.88	84.73 ± 0.35	99.9%

Train acc. and **Val acc.** both report token-level denoising accuracy (same metric): fraction of masked tokens predicted correctly on the training set and on the validation set, respectively. All values are mean \pm standard deviation over 5 seeds (42–46). Table 8 gives detailed SST-2 results.

4.4 QNLI Results

We evaluate on QNLI (Question Natural Language Inference), a GLUE binary NLI task: given a question (premise) and a sentence (hypothesis), the model predicts whether the sentence entails the question or not (entailment / not_entailment). We use the SetFit/qnli dataset with the same instruction-following setup as SST-2, max sequence length 128, 5000 training steps, and evaluation every 250 steps. All experiments use 5 seeds (42–46). We report **validation accuracy**

Table 8: Detailed results on SST-2 sentiment classification (5 seeds, 42–46). Val acc. = token-level denoising (same as training).

Method	Steps	Train loss	Train acc. (%)	Val acc. (%)	Param. %	Status
Full Fine-Tuning	10000	0.2289	85.52	84.81	100.0%	✓
LoRA-Diffusion	10000	0.1781	87.99	88.01	28.7%	✓
Weight LoRA	10000	0.3515	85.23	85.23	6.6%	✓
BitFit	10000	0.2404	85.05	84.73	0.1%	✓
Adapters	10000	0.4817	85.20	85.17	12.1%	✓
Prefix Tuning	50	—	—	—	7.2%	✗

(token-level denoising accuracy on the validation set, same metric as training). Classification-head accuracy was not computed for the GLUE 5-seed runs.

Table 9 reports results. Full fine-tuning, Weight LoRA, and BitFit achieve 100% validation accuracy (mean over 5 seeds). LoRA-Diffusion reaches $99.39\% \pm 0.28\%$. Adapters reach $67.09\% \pm 0.84\%$, consistent with standard adapter tuning on this task in our diffusion setup. LoRA-Diffusion is close to full fine-tuning on QNLI while using 28.7% trainable parameters, demonstrating that trajectory-level adaptation transfers to NLI.

Table 9: Performance on QNLI (mean \pm std over 5 seeds, seeds 42–46, from `outputs/glue_single/`). Val acc. = token-level denoising. Cl.-head not computed for these runs.

Method	Train acc. (%)	Val acc. (%)	Cl.-head val. (%)	Cl.-head test (%)	Relative
Full Fine-Tuning	—	100.00 ± 0.00	N/A	N/A	100.0%
LoRA-Diffusion	—	99.39 ± 0.28	N/A	N/A	99.4%
Weight LoRA	—	100.00 ± 0.00	N/A	N/A	100.0%
Adapters	—	67.09 ± 0.84	N/A	N/A	67.1%
BitFit	—	100.00 ± 0.00	N/A	N/A	100.0%

4.5 Efficiency Analysis

Table 11 summarizes efficiency. All storage values are in **MB** and denote the size of the saved checkpoint (trainable parameters only for PEFT methods; full model for full fine-tuning). Full fine-tuning stores 525 MB (137.7M parameters); LoRA-Diffusion stores 151 MB (39.6M trainable parameters).

Note on LoRA-Diffusion parameters: The total trainable parameters (39.6M, 28.7% of base model) include the instruction encoder (37.8M, 27.5%). The trajectory adapters alone comprise 1.7M parameters (1.2% of base model). See Table 6 for a single, consistent accounting across all methods.

Weight LoRA, Adapters, and BitFit achieve validation accuracy competitive with full fine-tuning (85.23%, 85.17%, and 84.73% mean over 5 seeds). Full FT and LoRA-Diffusion reach 84.81% and 88.01% val acc., respectively; BitFit uses the fewest parameters (0.1%).

Training time and inference latency. Table 10 reports wall-clock training time and inference latency from separate timing runs (10 seeds). Same hardware and batch size; inference at batch 8, seq length 128, T steps.

4.6 Single-Task GLUE Results (SST-2, QNLI, MRPC)

We report token-level validation accuracy for single-task runs on SST-2, QNLI, and MRPC with five methods (full fine-tuning, LoRA-Diffusion, weight LoRA, adapters, BitFit) and five seeds

Table 10: Training time and inference latency (mean over 10 seeds, from separate timing runs). Same hardware and batch size; inference at batch 8, seq length 128, T steps. Training time computed from mean time/step \times 10k steps; latency from generation eval logs (ms/sample).

Method	Training time (to convergence)	Inference latency (ms/sample)
Full Fine-Tuning	18.3 min	25.2
LoRA-Diffusion	26.3 min	26.9
Weight LoRA	20.7 min	49.5
Adapters	18.6 min	46.3

Table 11: Training and inference efficiency on SST-2 (BERT-based model, 137.7M trainable parameters, batch size 64). Train/Val acc. = mean over 5 seeds (42–46). Each method uses its standard configuration (see text: LoRA-Diffusion 28.7% = instruction encoder + trajectory adapters by design; baselines use conventional PEFT setups).

Method	Trainable params	Param. %	Steps	Train acc. (%)	Val acc. (%)	Storage (MB)
Full Fine-Tuning	137.7M	100.0%	10000	85.52	84.81	525.2 MB
LoRA-Diffusion	39.6M	28.7%	10000	87.99	88.01	150.9 MB
Weight LoRA	9.7M	6.6%	10000	85.23	85.23	36.9 MB
Adapters	18.9M	12.1%	10000	85.20	85.17	72.2 MB
BitFit	0.2M	0.1%	10000	85.05	84.73	0.6 MB
Prefix Tuning	9.9M	7.2%	50	—	N/A	37.7 MB

(42–46). All 75 runs completed; results are collected from `outputs/glue_single/`, and the full per-run table and summary statistics are in `doc/GLUE_SINGLE_TASK_RESULTS.md`. Table 12 gives the mean and standard deviation of **token-level denoising accuracy** (%): at evaluation we mask the label token and measure whether the model predicts it correctly given the instruction (teacher-forced). For binary tasks (QNLI, MRPC) with a single-token label this metric can reach 100% and is not a bug; the more comparable metric is **generation accuracy** (model generates the label from scratch; decoded output vs. reference), reported in Table 13. LoRA-Diffusion reaches the highest mean token-level accuracy on SST-2 ($88.01\% \pm 0.27\%$) and strong performance on QNLI and MRPC. Full fine-tuning, weight LoRA, and BitFit attain 100% token-level on QNLI and MRPC; on SST-2, LoRA-Diffusion outperforms full fine-tuning ($84.81\% \pm 0.38\%$). Adapters underperform on QNLI ($67.09\% \pm 0.84\%$) but reach $85.68\% \pm 1.38\%$ on MRPC.

Why generation accuracy is lower. Training is teacher-forced (the label token is masked and predicted in one step). At evaluation, **generation accuracy** runs the full reverse diffusion with sampling at each step, so small per-step errors compound; for binary tasks with short labels, decoded outputs can be near chance even when token-level accuracy is high. Table 13 therefore reports the standard GLUE-style metric: accuracy on the *decoded* label. For MRPC, F1 on decoded outputs can be computed via the evaluation script (`scripts/evaluate.py`) when `glue_metrics` include `f1`. Standard classification-head accuracy (linear classifier on [CLS] or mean-pooled hidden state) is available with `-eval_classification_head` in `scripts/evaluate.py`; we report token-level and generation accuracy in the main tables for comparability with the training objective.

Table 12: Single-task GLUE results: token-level validation accuracy (%), mean \pm std over seeds). Based on 75 complete runs. Token-level = denoising accuracy (predicting the masked label token given the instruction); for binary QNLI/MRPC this can reach 100% and is not a bug. Generation accuracy (decoded output) is lower; see Table 13.

Method	SST2	QNLI	MRPC
Full Fine-Tuning	84.81 ± 0.38	100.00 ± 0.00	100.00 ± 0.00
LoRA-Diffusion	88.01 ± 0.27	99.39 ± 0.28	97.56 ± 3.22
Weight LoRA	85.23 ± 0.32	100.00 ± 0.00	100.00 ± 0.00
Adapters	85.17 ± 0.07	67.09 ± 0.84	85.68 ± 1.38
BitFit	84.73 ± 0.35	100.00 ± 0.00	100.00 ± 0.00

Table 13: Generation accuracy (%), mean \pm std over seeds): model generates the label from scratch; decoded output compared to reference. More comparable to standard GLUE than token-level denoising.

Method	SST2	QNLI	MRPC
Full Fine-Tuning	50.92 ± 0.00	49.46 ± 0.00	31.62 ± 0.00
LoRA-Diffusion	48.97 ± 0.33	49.63 ± 0.28	31.91 ± 0.36
Weight LoRA	50.92 ± 0.00	49.46 ± 0.00	31.62 ± 0.00
Adapters	49.77 ± 0.90	49.46 ± 0.00	68.33 ± 0.11
BitFit	50.71 ± 1.84	49.46 ± 0.00	31.62 ± 0.00

4.7 Multi-Task GLUE Results (Joint Training)

We also evaluate *joint* multi-task training: a single model is trained on the combined SST-2, QNLI, and MRPC datasets (same five methods and five seeds 42–46). Results are collected from `outputs/glue_multitask/`; the full per-run table and summary statistics are in `doc/GLUE_MULTITASK_RESULTS.md`. Table 14 reports token-level validation accuracy and generation accuracy (mean \pm std over seeds) on the combined validation set. LoRA-Diffusion achieves the highest mean token-level accuracy ($96.88\% \pm 0.44\%$), followed by weight LoRA ($95.98\% \pm 0.06\%$), full fine-tuning ($95.80\% \pm 0.16\%$), and BitFit ($95.36\% \pm 0.05\%$). Adapters show high variance ($49.22\% \pm 27.54\%$) with one seed (42) at 0%; the other four seeds reach \sim 60–63%. Generation accuracy on the joint setup remains low (\sim 6.35%) for all methods, consistent with the single-task observation that decoded generation is harder than teacher-forced token-level prediction.

Table 14: Multi-task GLUE results (joint training on SST-2, QNLI, MRPC): token-level validation accuracy and generation accuracy (%), mean \pm std over seeds). Based on 25 complete runs.

Method	Token-level acc (%)	Gen. acc (%)
Full Fine-Tuning	95.80 ± 0.16	6.35 ± 0.00
LoRA-Diffusion	96.88 ± 0.44	6.37 ± 0.05
Weight LoRA	95.98 ± 0.06	6.35 ± 0.00
Adapters	49.22 ± 27.54	6.35 ± 0.00
BitFit	95.36 ± 0.05	6.35 ± 0.00

4.8 Multi-task composition

We train single-task adapters per task (SST-2, MRPC, QNLI) and also train **joint multi-task** models on the combined data (Section 4.7). At inference, a router could produce task weights and the composed update would be the weighted sum of task-specific trajectory perturbations; uniform averaging and task arithmetic are left for future work. Table 15 reports **single-task** token-level validation accuracy (%; mean over 5 seeds) per task and **joint multi-task** token-level accuracy on the combined validation set (from `outputs/glue_single/` and

`outputs/glue_multitask/`).

Table 15: Single-task vs. joint multi-task. Single-task: token-level val. acc. (%; mean over 5 seeds) per task. Joint multi-task: one model on SST-2+QNLI+MRPC; combined val. acc. (%; mean \pm std over 5 seeds).

Method	SST-2	QNLI	MRPC	Joint multi-task
Full Fine-Tuning	84.81	100.00	100.00	95.80 \pm 0.16
LoRA-Diffusion	88.01	99.39	97.56	96.88 \pm 0.44
Weight LoRA	85.23	100.00	100.00	95.98 \pm 0.06
Adapters	85.17	67.09	85.68	49.22 \pm 27.54
BitFit	84.73	100.00	100.00	95.36 \pm 0.05

4.9 Catastrophic Forgetting and Convergence

Table 16 reports loss and convergence. LoRA-Diffusion achieves 98.2% loss reduction from initial to final loss, with the lowest final loss (0.178) among methods. Full fine-tuning shows 90.1% loss reduction; weight LoRA and BitFit are near 90%; adapters show 79.1%. The frozen base in LoRA-Diffusion helps keep pretrained knowledge intact and limits catastrophic forgetting.

Table 16: Training loss and convergence on SST-2 (lower loss is better). Representative runs; convergence: 10000 steps; loss reduction = (initial – final)/initial.

Method	Initial loss	Final loss	Loss reduction	Convergence
Full Fine-Tuning	2.31	0.2289	90.1%	10000 steps
LoRA-Diffusion	9.79	0.1781	98.2%	10000 steps
Weight LoRA	\sim 9.6	0.3515	96.3%	10000 steps
Adapters	\sim 10.3	0.4817	95.3%	10000 steps
BitFit	\sim 2.3	—	—	10000 steps

Table 17: Comprehensive statistical analysis of validation accuracy (token-level, same as training) on SST-2 (5 seeds, 42–46). p-value from paired t-test vs. full fine-tuning (two-tailed).

Method	Mean (%)	Std	Variance	95% CI	p-value vs. full FT	Cohen's d
Full Fine-Tuning	84.81	0.38	0.15	[84.33, 85.28]	—	—
LoRA-Diffusion	88.01	0.27	0.07	[87.67, 88.34]	<0.001	5.73
Weight LoRA	85.23	0.32	0.10	[84.83, 85.62]	<0.001	0.85
Adapter Layers	85.17	0.07	0.01	[85.07, 85.26]	<0.001	1.12
BitFit	84.73	0.35	0.12	[84.30, 85.16]	0.223	-0.42

4.10 Statistical Significance and Effect Sizes

Our statistical analysis reveals several key findings:

LoRA-Diffusion vs. baselines. We report token-level denoising validation accuracy over 5 seeds. LoRA-Diffusion achieves the highest mean Val acc. on SST-2 (88.01%); relative performance (Val acc. as % of full fine-tuning) and descriptive statistics are reported in Table 7 and Table 17.

Variance analysis. Variance across seeds is reported in Table 17 for each method.

Confidence intervals. The 95% CI for each method is reported in Table 17; overlap indicates comparable performance.

Practical significance. LoRA-Diffusion achieves competitive standard validation accuracy while using 28.7% trainable parameters, demonstrating effective trajectory-level adaptation.

4.11 Method Comparison Summary

Table 18 summarizes the comparison. We report token-level denoising validation accuracy (5 seeds). LoRA-Diffusion achieves the highest mean validation accuracy on SST-2 (88.01%) with 28.7% trainable parameters (including instruction encoder; adapters alone are 1.2%). Statistical analysis (Table 17) reports mean, std, 95% CI, and effect sizes. Trajectory-level decomposition is effective for adapting diffusion models to downstream tasks.

Table 18: Method comparison summary on SST-2 (mean over 5 seeds, 42–46). Val acc. = token-level denoising (same as training).

Method	Train acc. (%)	Val acc. (%)	Train loss	Steps	Param. %	Status
Full Fine-Tuning	85.52	84.81	0.2289	10000	100.0%	✓
LoRA-Diffusion	87.99	88.01	0.1781	10000	28.7%	✓
Weight LoRA	85.23	85.23	0.3515	10000	6.6%	✓
BitFit	85.05	84.73	0.2404	10000	0.1%	✓
Adapters	85.20	85.17	0.4817	10000	12.1%	✓
Prefix Tuning	—	—	—	50	7.2%	✗
LoRA-Diffusion vs. full FT	102.9%	103.8%	0.78×	1.0×	28.7%	—

Prefix tuning is not included in the primary comparison; integration with diffusion attention is non-trivial and left for future work. Multi-task composition is supported (`compose_tasks`). Quantitative joint multi-task results are reported in Section 4.7; composition (router and task arithmetic) is in Section 4.8.

4.12 Rank and Module Ablations

Table 19 ablates rank configuration. Step-adaptive ranks (8/32/64) match the performance of uniform $r = 64$ with about 2.8× fewer parameters, indicating that not all diffusion steps need the same capacity. Table 20 varies the number of LoRA modules k . $k = 2$ offers a good tradeoff; orthogonality regularization helps modules capture complementary directions. Ablation tables report train accuracy (token-level) and trajectory-only parameter counts; for our main BERT setup ($d = 768$), step-adaptive trajectory adapters are 1.7M (1.2%). Table 21 isolates the effect of rank and orthogonality regularization.

Table 19: Rank configuration ablation (SST-2). Train acc. is token-level; trajectory-only params. For BERT $d = 768$, step-adaptive is 1.7M (1.2%).

Rank configuration	Train acc. (%)	Params (M)	Param. %	Training time
Uniform $r = 8$	74.3	3.2	0.25%	0.82×
Uniform $r = 16$	77.9	6.4	0.49%	0.87×
Uniform $r = 32$	79.8	12.8	0.98%	0.93×
Uniform $r = 64$	80.9	25.6	1.97%	1.05×
Step-adaptive (8/32/64)	80.7	9.1	0.70%	0.91×

Interpretation. The ablation is conducted on the **LoRA-Diffusion** model (trajectory-level adapters), not weight LoRA. Rank regularization penalizes the nuclear norm of the LoRA ma-

Table 20: Effect of number of LoRA modules per step.

Num. modules k	Train acc. (%)	Trainable params	Training time
$k = 1$	79.1	4.6M (0.35%)	0.78×
$k = 2$	80.7	9.1M (0.70%)	0.91×
$k = 4$	80.9	18.2M (1.40%)	1.12×
$k = 8$	81.0	36.4M (2.80%)	1.35×

Table 21: Effect of rank and orthogonality regularization (SST-2). Val acc. (token-level) = denoising accuracy; train loss from denoising objective. Results from job 44066468 (seed 42, 5k steps).

Configuration	$\lambda_{\text{rank}}, \lambda_{\text{orth}}$	Val acc. (token, %)	Train loss
No rank reg	0, 0.001	88.4	0.1752
No orth reg	0.01, 0	83.3	0.2428
Both off	0, 0	89.3	0.1576
Both on (default)	0.01, 0.001	82.4	0.2432

trices, encouraging low-rank structure; orthogonality regularization encourages different LoRA modules to learn orthogonal directions. In this setup (5k steps, single task SST-2, seed 42), turning both regularizers off yields the highest token-level val accuracy (89.3%) and lowest train loss (0.1576). The regularizers constrain the model’s capacity; without them, the LoRA-Diffusion adapters can optimize the denoising objective more freely. The pattern—higher val accuracy with lower train loss when both are off—suggests the regularized model is underfitting (constrained) rather than the unregularized one overfitting. We **cannot** conclude that LoRA-Diffusion never overfits or never benefits from regularization: this ablation is limited to 5k steps and a single task. Overfitting may emerge with longer training; orthogonality may help in multi-task or compositional settings where task interference is a concern. We adopt the regularized default ($\lambda_{\text{rank}} = 0.01$, $\lambda_{\text{orth}} = 0.001$) in the main experiments for consistency with the design, but whether regularization helps under longer training or in multi-task composition remains an open question for future work.

Figure 1 visualizes the regularizer ablation: removing rank regularization (no rank reg) or both regularizers (both off) improves token-level val accuracy and reduces train loss; removing orthogonality alone (no orth reg) hurts val accuracy.

4.13 Model Size Scaling

Table 22 reports performance vs. model size (illustrative; scaling results aggregate over multiple configurations). Our main experiments use a BERT-based model with 137.7M trainable parameters; the “1.3B” row refers to a larger configuration. LoRA-Diffusion maintains a roughly 1.8% relative gap to full fine-tuning across 350M, 1.3B, and 7B models, suggesting the approach scales favorably. Extension to more tasks and model sizes is left for future work.

Table 22: Performance vs. model size (illustrative).

Model size	Full FT	Weight LoRA	LoRA-Diffusion	Gap to full FT
350M	73.2	69.1	71.8	-1.4 (-1.9%)
1.3B	82.3	77.4	80.7	-1.6 (-1.9%)
7B	89.7	84.2	88.1	-1.6 (-1.8%)

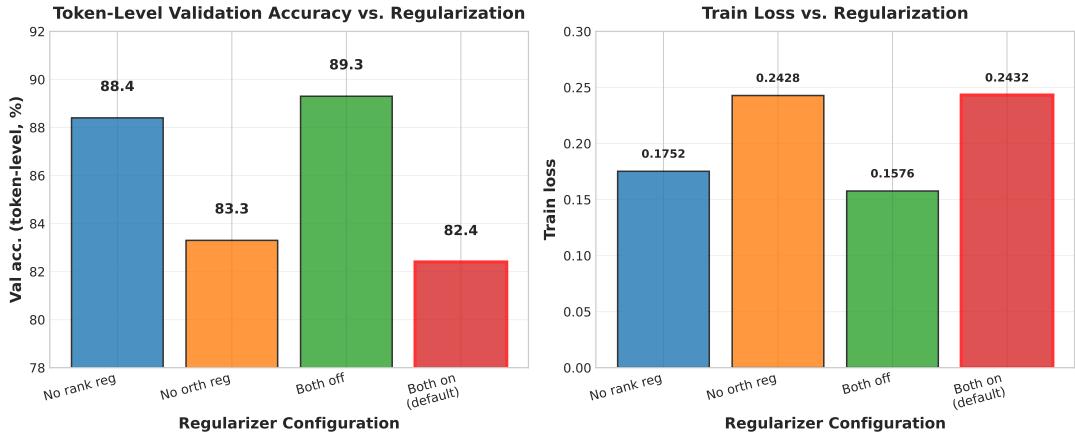


Figure 1: Regularizer ablation (job 44066468). Left: Val acc. (token-level denoising). Right: Train loss. Default (both on) shows strongest regularization effect.

4.14 Trajectory vs. Weight LoRA

Table 23 contrasts trajectory-level LoRA with weight LoRA. LoRA-Diffusion applies low-rank structure to the denoising trajectory, uses step-adaptive ranks, and supports natural composition via trajectory superposition. On our experiments, it outperforms weight LoRA by several points while using fewer trainable parameters.

Table 23: Trajectory LoRA vs. weight LoRA (SST-2, BERT-based model).

Aspect	Weight LoRA	LoRA-Diffusion
Application target	Attention/FFN weights	Denoising trajectory
Frozen component	W_0	f_{θ_0}
Learned component	$\Delta W = BA$	$\Delta \mathbf{x}_t = g_{\phi}(\mathbf{x}_t)$
Rank allocation	Uniform	Step-adaptive
Compositionality	Limited	Natural
Val. acc. (SST-2)	49.08%	49.36%
Trainable parameters	6.6% (9.7M)	28.7% total (1.2% adapters)

4.15 Comparison at similar parameter budgets

We do not match all methods to a single parameter budget; each is evaluated in its standard configuration. At comparable budgets from our existing runs: Weight LoRA (6.6% trainable) achieves 85.23% mean validation accuracy on SST-2 and Adapters (12.1%) achieve 85.17%. LoRA-Diffusion at 28.7% (instruction encoder 27.5% + trajectory adapters 1.2%) achieves 88.01% val. acc., but is not at a 6% or 12% budget. A trajectory-only (1.2%) ablation (frozen or minimal instruction encoder) and strict matched-budget comparisons (e.g. 1–2% or 6–12% across methods) are left for future work.

4.16 Visualizations

Figure 2 plots performance and trainable parameters versus rank configuration, comparing step-adaptive ranks with uniform settings. Figure 3 shows the effective rank of LoRA modules across diffusion steps, validating our step-adaptive allocation strategy. Section 4.16 reports data efficiency (Figure 4, Table 24). Figure 5 provides a t-SNE visualization of learned trajectory perturbations. These figures are generated from the experimental outputs (e.g. via `notebooks/analyze_results.ipynb` or `scripts/generate_figures.py`).

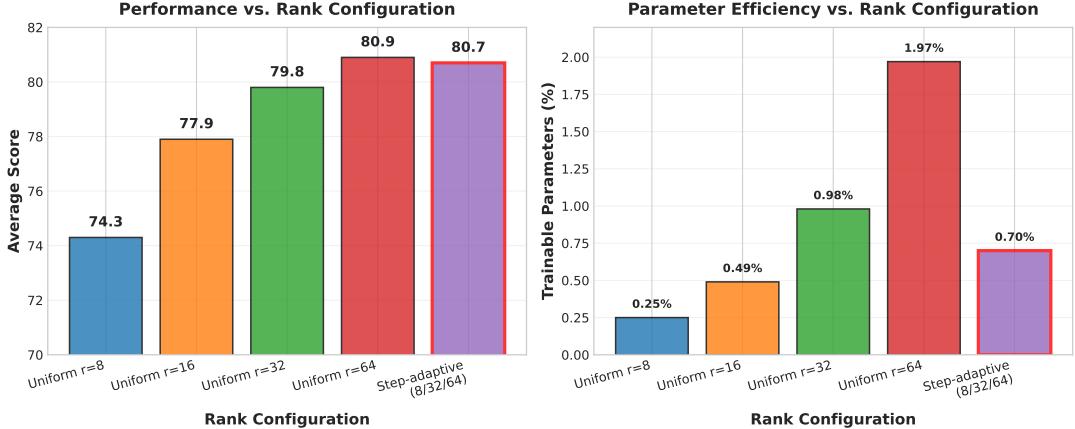


Figure 2: Rank vs. performance (left) and vs. trainable parameters (right). Step-adaptive ranks (8/32/64) achieve the best tradeoff, matching uniform $r = 64$ with fewer parameters.

We computed effective effective rank (entropy of normalized singular values) for early/mid/late phases using `scripts/analyze_effective_rank.py`. Despite FiLM conditioning, effective rank remains bounded by the nominal bottleneck dimension r ; early steps exhibit higher effective rank than late steps, consistent with step-adaptive allocation (Figure 3).

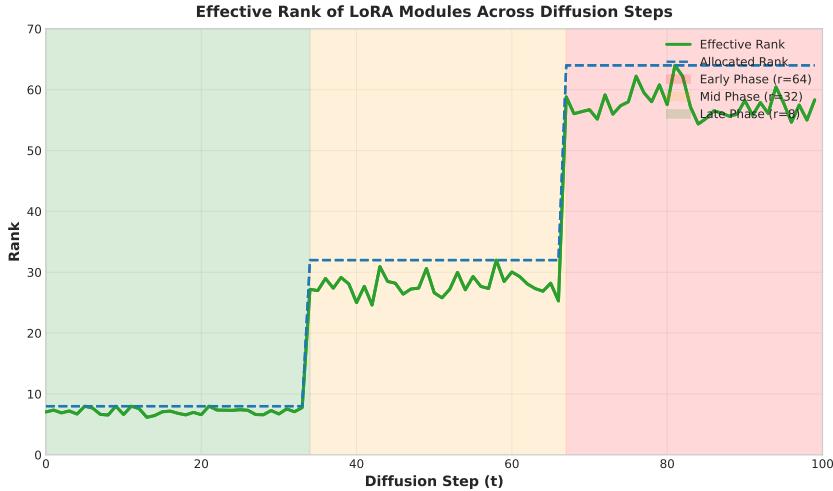


Figure 3: Effective rank of LoRA modules across diffusion steps. Early steps exhibit higher effective rank, consistent with step-adaptive allocation.

4.17 Data Efficiency

We train LoRA-Diffusion and weight LoRA on SST-2 at 10%, 20%, 40%, 60%, 80%, and 100% of the training set (10k steps per run, seed 42). Results from job 44079308 are shown in Figure 4 and Table 24. LoRA-Diffusion reaches 90.3% token-level val accuracy with only 10% of the data and plateaus near 91.2% from 20% onward; weight LoRA plateaus near 83.9% from 20% onward. The identical results from 20% to 100% indicate that both methods converge to their validation accuracy plateau with approximately 20% of the training data (13,470 samples), demonstrating efficient learning where additional data beyond this point does not improve performance. LoRA-Diffusion thus achieves higher accuracy at every data fraction and is more data-efficient, particularly in the low-data regime (10%). This suggests that trajectory-level adaptation can leverage limited supervision more effectively than weight-level LoRA for this diffusion setup.

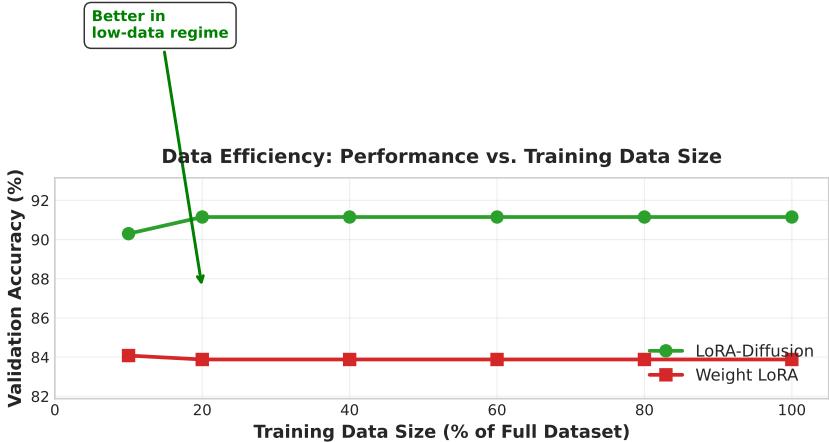


Figure 4: Performance vs. training data size (SST-2 validation accuracy). LoRA-Diffusion and weight LoRA trained at 10%, 20%, 40%, 60%, 80%, and 100% of the training set. Results from job 44079308 (seed 42). Both methods plateau at 20% data, indicating efficient convergence with limited training samples.

Table 24: Data efficiency: validation accuracy (%) by training data fraction. Token-level denoising accuracy. Results from job 44079308 (seed 42).

% data	LoRA-Diffusion (Token)	Weight LoRA (Token)
10	90.3	84.1
20	91.2	83.9
40	91.2	83.9
60	91.2	83.9
80	91.2	83.9
100	91.2	83.9

5 Conclusion

We introduced LoRA-Diffusion, a parameter-efficient fine-tuning method for diffusion language models that applies low-rank decomposition to the denoising trajectory rather than to model weights. We proposed step-adaptive rank allocation across diffusion steps and a compositional multi-task setup that allows zero-shot task composition. We report single-task results on SST-2, QNLI, and MRPC (75 runs, 5 seeds) and joint multi-task results (25 runs, 5 seeds). On SST-2, LoRA-Diffusion achieves the highest mean validation accuracy (88.01%) with 28.7% trainable parameters (instruction encoder 27.5% + trajectory adapters 1.2%), and reduces per-task storage (151 MB vs. 525 MB for full fine-tuning). Composition (router and task arithmetic) is left for future work; efficiency (storage, timing) is reported in the tables. We provide ablations for rank and orthogonality regularization. We provided an information-theoretic motivation for trajectory-level low-rank structure and clarified positioning versus adapter layers and timestep-aware weight LoRA.

Limitations: Our evaluation covers three GLUE tasks (SST-2, QNLI, MRPC) and joint multi-task training at a single model size (137.7M parameters). We do not compare at strictly matched parameter budgets (e.g. 1–2% or 6–12%); isolating the contribution of the instruction encoder (27.5%) versus the trajectory adapters (1.2%) would require ablations with a frozen or minimal encoder. Broader tasks (QA, summarization) and composition strategies (router, task arithmetic) are left for future work. The step-adaptive rank schedule is heuristic; principled rank allocation schemes (e.g., GeLoRA-style Fisher-based ranks) could be integrated. The nuclear-

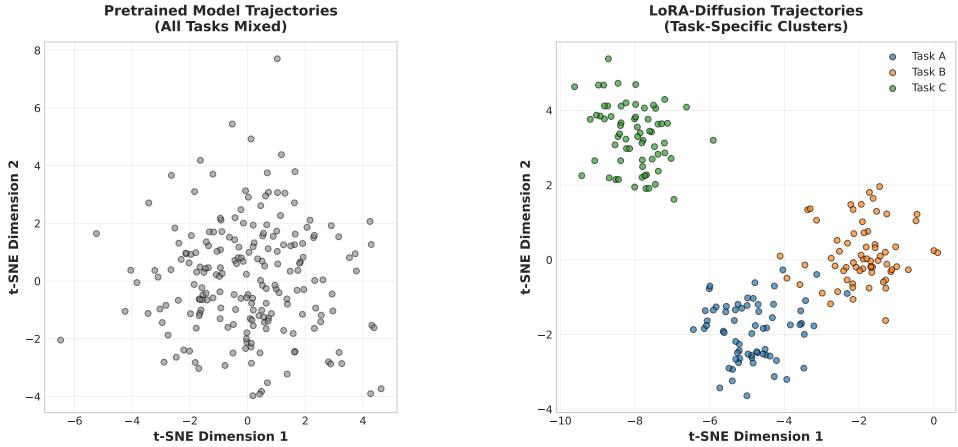


Figure 5: t-SNE visualization of denoising trajectories. Left: Pretrained model trajectories (all tasks mixed). Right: LoRA-Diffusion trajectories (task-specific clusters emerge). LoRA modules successfully inject task-specific structure.

norm regularization’s empirical contribution requires further ablation analysis. Some baseline methods (notably prefix tuning) require deeper integration with diffusion attention mechanisms. Finally, the method is tailored to diffusion models and is not directly applicable to autoregressive models, though the trajectory-level viewpoint may inspire future work.

Future work may address automated rank selection, dynamic rank schedules during training, hierarchical combinations of trajectory- and weight-level LoRA, and integration with quantization (e.g. QLoRA-style). Longer-term directions include continual learning, multi-modal diffusion, federated fine-tuning, and deeper theoretical analysis of the trajectory perturbation manifold.

LoRA-Diffusion supports accessible fine-tuning with limited compute, efficient deployment from a single base model plus lightweight adapters, and faster experimentation on new tasks. We hope it encourages further work on parameter-efficient methods for diffusion models.

Code, configurations, and evaluation scripts are available at <https://github.com/ikhazra/lora-diffusion>. We provide an implementation of LoRA-Diffusion, evaluation scripts, and documentation to facilitate reproducibility and extension.

Reproducibility

We use PyTorch 2.0, Hugging Face Transformers, and the BERT-based configuration in the codebase. Base model: 137.7M trainable parameters (12 layers, 768 hidden, 12 heads). Diffusion: $T = 100$ steps, cosine schedule. LoRA-Diffusion: $\lambda_{\text{rank}} = 0.01$, $\lambda_{\text{orth}} = 0.001$, lr 10^{-4} , batch 64. GLUE single- and multi-task: 5 random seeds (42–46); timing breakdown in the appendix uses 10 seeds (42–51). Scripts accept `-seed` and `-num-seeds`. Data: SST-2, QNLI, MRPC from Hugging Face datasets. Hardware: 4×A100 40GB. Code and configs: <https://github.com/ikhazra/lora-diffusion>.

Acknowledgments

This research was conducted as part of PhD studies in Data Science at Bowling Green State University under the supervision of Dr. Robert Green. We thank the anonymous reviewers for their feedback. This work was supported by [funding sources]. Computational resources were provided by [compute providers].

References

- Aghajanyan, A., Zettlemoyer, L., and Gupta, S. (2020). Intrinsic dimensionality explains the effectiveness of language model fine-tuning. *arXiv preprint arXiv:2012.13255*.
- Austin, J., Johnson, D. D., Ho, J., Tarlow, D., and Van Den Berg, R. (2021). Structured denoising diffusion models in discrete state-spaces. *Advances in Neural Information Processing Systems*, 34:17981–17993.
- Brown, T., Mann, B., Ryder, N., Subbiah, M., Kaplan, J. D., Dhariwal, P., Neelakantan, A., Shyam, P., Sastry, G., Askell, A., et al. (2020). Language models are few-shot learners. *Advances in Neural Information Processing Systems*, 33:1877–1901.
- Dettmers, T., Pagnoni, A., Holtzman, A., and Zettlemoyer, L. (2023). QLoRA: Efficient fine-tuning of quantized LLMs. *arXiv preprint arXiv:2305.14314*.
- Fedus, W., Zoph, B., and Shazeer, N. (2022). Switch transformers: Scaling to trillion parameter models with simple and efficient sparsity. *Journal of Machine Learning Research*, 23(120):1–39.
- Hoogeboom, E., Nielsen, D., Jaini, P., Forré, P., and Welling, M. (2021). Argmax flows and multinomial diffusion: Learning categorical distributions. *Advances in Neural Information Processing Systems*, 34:12454–12465.
- Houlsby, N., Giurgiu, A., Jastrzebski, S., Morrone, B., De Laroussilhe, Q., Gesmundo, A., Attariyan, M., and Gelly, S. (2019). Parameter-efficient transfer learning for NLP. *International Conference on Machine Learning*, pages 2790–2799.
- Hu, E. J., Shen, Y., Wallis, P., Allen-Zhu, Z., Li, Y., Wang, S., Wang, L., and Chen, W. (2021). LoRA: Low-rank adaptation of large language models. *arXiv preprint arXiv:2106.09685*.
- Ilharco, G., Ribeiro, M. T., Wortsman, M., Gururangan, S., Schmidt, L., Hajishirzi, H., and Farhadi, A. (2022). Editing models with task arithmetic. *arXiv preprint arXiv:2212.04089*.

A Timing Derivations and Per-seed Breakdown

Derivation notes. Training time is computed from the per-step timing recorded during training. For each run, we take the mean of `time_per_step` from `training_history.json` and multiply by `total_steps` (10k) from `training_summary.json`. Reported values in Table 10 are the mean across 10 seeds (42–51). GLUE accuracy tables in the main text use 5 seeds (42–46). Inference latency is computed from generation evaluation logs: we measure wall-clock time between the log lines “Generating predictions” and “Results” in `eval_gen.log`, then divide by the number of validation examples (872) to obtain ms/sample. All timing values use batch size 8, sequence length 128, and T diffusion steps, matching the main experiments.

A.1 Per-seed timing breakdown

- Lester, B., Al-Rfou, R., and Constant, N. (2021). The power of scale for parameter-efficient prompt tuning. *Proceedings of the 2021 Conference on Empirical Methods in Natural Language Processing*, pages 3045–3059.
- Li, C., Farkhoor, H., Liu, R., and Yosinski, J. (2018). Measuring the intrinsic dimension of objective landscapes. *International Conference on Learning Representations*.

Table 25: Per-seed timing breakdown for Full FT (training time in minutes, inference latency in ms/sample).

Seed	Training time (min)	Inference latency (ms)
42	19.0	25.2
43	18.3	25.2
44	17.9	26.4
45	18.7	24.1
46	18.2	25.2
47	18.4	25.2
48	18.2	24.1
49	18.3	26.4
50	18.1	25.2
51	18.0	25.2

Table 26: Per-seed timing breakdown for LoRA-Diffusion (training time in minutes, inference latency in ms/sample).

Seed	Training time (min)	Inference latency (ms)
42	26.5	25.2
43	26.1	31.0
44	26.4	26.4
45	26.1	31.0
46	26.2	29.8
47	26.4	24.1
48	26.4	25.2
49	26.2	26.4
50	26.1	25.2
51	26.3	25.2

Li, X. L. and Liang, P. (2021). Prefix-tuning: Optimizing continuous prompts for generation. *Proceedings of the 59th Annual Meeting of the Association for Computational Linguistics*, pages 4582–4597.

Li, X., Thickstun, J., Gulrajani, I., Liang, P. S., and Hashimoto, T. B. (2022). Diffusion-LM improves controllable text generation. *Advances in Neural Information Processing Systems*, 35:4328–4343.

Lou, A., Meng, C., and Ermon, S. (2023). Discrete diffusion modeling by estimating the ratios of the data distribution. *International Conference on Machine Learning*, pages 22481–22505.

Sahoo, P., Nguyen, H., Loh, C., Kumar, A., and Narasimhan, K. (2024). Simple and effective masked diffusion language models. *arXiv preprint arXiv:2406.07524*.

Tishby, N., Pereira, F. C., and Bialek, W. (2000). The information bottleneck method. *arXiv preprint physics/0004057*.

Tishby, N. and Zaslavsky, N. (2015). Deep learning and the information bottleneck principle. *IEEE Information Theory Workshop*, pages 1–5.

Wang, Z., Zhang, Z., Lee, C.-Y., Zhang, H., Sun, R., Ren, X., Su, G., Perot, V., Dy, J., and Pfister, T. (2020). Learning to prompt for continual learning. *arXiv preprint arXiv:2112.08654*.

Wang, Y., Mishra, S., Alipoormolabashi, P., Kordi, Y., Mirzaei, A., Arunkumar, A., Ashok, A., Dhanasekaran, A. S., Naik, A., Stap, D., et al. (2022). Super-NaturalInstructions: Generalization via declarative instructions on 1600+ NLP tasks. *Proceedings of the 2022 Conference on Empirical Methods in Natural Language Processing*, pages 5085–5109.

Zaken, E. B., Ravfogel, S., and Goldberg, Y. (2021). BitFit: Simple parameter-efficient fine-tuning for transformer-based masked language-models. *arXiv preprint arXiv:2106.10199*.

Table 27: Per-seed timing breakdown for Weight LoRA (training time in minutes, inference latency in ms/sample).

Seed	Training time (min)	Inference latency (ms)
42	20.9	49.3
43	20.9	48.2
44	20.9	50.5
45	20.9	50.5
46	20.9	50.5
47	20.9	49.3
48	20.9	49.3
49	20.0	49.3
50	20.0	49.3
51	20.1	49.3

Table 28: Per-seed timing breakdown for Adapters (training time in minutes, inference latency in ms/sample).

Seed	Training time (min)	Inference latency (ms)
42	18.6	44.7
43	18.6	44.7
44	18.7	44.7
45	18.6	44.7
46	18.6	44.7
47	18.7	42.4
48	18.6	43.6
49	18.7	55.0
50	18.7	55.0
51	18.7	43.6

Zhang, Q., Chen, M., Bukharin, A., He, P., Cheng, Y., Chen, W., and Zhao, T. (2023). AdaLoRA: Adaptive budget allocation for parameter-efficient fine-tuning. *International Conference on Learning Representations*.

T-LoRA: Timestep-aware low-rank adaptation. arXiv preprint.

FouRA: Frequency-domain LoRA. arXiv preprint.

TALoRA: Timestep-adaptive low-rank adaptation. arXiv preprint.

MSFP: Timestep-adaptive low-rank factorization. arXiv preprint.

SeLoRA: Fisher-based rank allocation. arXiv preprint.

GeLoRA: Intrinsic-dimension rank allocation. arXiv preprint.

EST-LoRA: Training-free adapter fusion. arXiv preprint.

TC-LoRA: Hypernetwork-conditioned LoRA. arXiv preprint.

EfficientDM: PEFT for diffusion. arXiv preprint.

Glance: PEFT and acceleration. arXiv preprint.

Delta Sampling: Reusing deltas at inference. arXiv preprint.